



Figure 17: NIAH benchmark for LLaMA-3-8B-Instruct with KV cache size=128 under 8k context length