



Figure 18: NIAH benchmark for Mistral-7B-Instruct with KV cache size=128 under 32k context length